

Artificial General Intelligence (AGI): Bridging the Gap Between AI and Humans

¹Pranav Verma, ²Varun Arora

¹Student, ²Student

¹Lotus Valley International School, Noida, India

¹pranav@verma.net.in, ²varunchamp26@gmail.com

Abstract- Artificial Intelligence (AI) has undergone remarkable evolution, transitioning from theoretical constructs to practical applications that significantly impact various sectors of society. Despite these advancements, the pursuit of Artificial General Intelligence (AGI)—machines capable of performing any intellectual task akin to human capabilities—remains a formidable challenge. This paper delves into the historical development of AI, analyzing its progression from early concepts to contemporary implementations, and highlights the current capabilities and limitations of narrow AI systems. We explore the multifaceted challenges and opportunities presented by AGI, particularly focusing on its potential benefits and associated risks. As AGI systems become increasingly integrated into critical areas such as healthcare and finance, understanding their implications becomes essential. Ethical considerations and societal impacts are examined in depth, emphasizing the need for responsible AI development practices. The paper concludes by offering insights into pathways for achieving AGI while underscoring the importance of aligning technological advancements with fundamental human values. This revised abstract provides a clearer overview of your paper's content, emphasizing the significance of AGI while maintaining a focus on historical context, current challenges, and ethical considerations.

I. Introduction

Artificial Intelligence (AI) has transformed from a conceptual framework into a powerful tool that influences multiple aspects of daily life, from healthcare to finance and entertainment. Narrow AI systems excel at specific tasks—such as image recognition, natural language processing, and game playing—but lack the ability to generalize knowledge across diverse domains. The goal of Artificial General Intelligence (AGI) is to create machines capable of understanding, learning, and applying knowledge in a manner akin to human intelligence. This paper discusses AGI within the broader context of AI research, focusing on technical challenges, current advancements, ethical implications, and future directions as we strive toward realizing AGI.

II. Evolution of Artificial Intelligence

AI has its roots in the 20th century, with early mathematicians such as Alan Turing laying the groundwork with concepts like the Turing Test, which measures a machine's ability to exhibit human-like behavior. The formal birth of AI occurred at the 1956 Dartmouth Conference, where researchers convened to explore the possibility of machines trying to mimic human intelligence [1]. Early AI systems were based on symbolic reasoning and rule-based approaches. However, these systems proved limited in real-world applications. The rise of machine learning shifted the paradigm by enabling machines to learn from data rather than relying on explicit programming. In particular, deep learning—a subfield of machine learning that uses multi-layered neural networks—led to significant advancements in image recognition, natural language processing, and game playing (e.g., AlphaGo, developed by DeepMind). Today's AI systems are powerful but limited to narrow tasks. While they can outperform humans in specific areas, such as image classification or strategic games, they lack the ability to generalize across different contexts. Narrow AI systems, like chatbots or recommendation engines, excel within their domain but cannot handle tasks outside their programming.

III. Understanding Artificial General Intelligence (AGI)

A. What is AGI?

Artificial General Intelligence (AGI) refers to an AI system with the ability to understand, learn, and apply knowledge across a variety of domains, much like human intelligence. Unlike narrow AI, AGI would not require retraining to perform tasks outside of its original programming. AGI aims to:

- Generalize knowledge across different fields.
- Reason abstractly and perform complex decision-making.
- Learn continuously and adapt to new situations.

AGI would have the potential to disrupt industries, revolutionize scientific research, and transform how humans interact with technology [2].

Aspect	Narrow AI	AGI
Task Scope	Designed for specific applications such as autonomous driving systems (e.g., Tesla Autopilot) and diagnostic tools in radiology (e.g., AI algorithms for detecting tumors in mammograms).	Capable of performing diverse tasks across various domains, such as natural language understanding, complex problem-solving in scientific research, and adaptive learning in dynamic environments.
Learning	Relies on extensive labeled datasets for training; for example, models like BERT require millions of text samples to understand language context.	Employs advanced learning paradigms such as few-shot learning and transfer learning, allowing it to adapt quickly to new tasks with minimal examples, similar to human cognitive flexibility.
Adaptability	Limited to the domain it was trained on; for instance, an AI trained for medical imaging cannot apply its knowledge to financial forecasting.	Exhibits robust cross-domain adaptability, allowing it to generalize knowledge from one area (e.g., language processing) to another (e.g., visual recognition), akin to human reasoning processes.
Human-Like Reasoning	Operates based on predefined algorithms without true understanding; for example, rule-based systems in chatbots often fail in ambiguous situations due to lack of contextual awareness.	Capable of abstract reasoning and complex decision-making; can generate novel solutions by synthesizing information from various fields, enhancing its ability to tackle intricate problems.

AGI and narrow AI differ fundamentally in their capabilities and applications, with narrow AI being limited to specific tasks while AGI aims for broader cognitive functions [3].

C. World Models and Viability in AGI

World Models and Reasoning

In the field of AGI, a "world model" encompasses the entirety of an agent’s knowledge, including skills and knowledge about the world. The development of intelligence, therefore, becomes synchronous with the expansion and refinement of this model. World models can vary significantly in complexity. A simple world model might only account for simple needs and actions, such as a basic survival mechanism. In contrast, a complex world model, like that of humans, considers past experiences, anticipates future outcomes, and represents numerous states of the world. Knowledge within a world model originates from three primary sources:

- **Prior Knowledge:** Initial assumptions and information given to the agent as a foundation.
- **Perceived Knowledge:** Information gathered in real-time through sensory inputs and direct observation.
- **Derived Knowledge:** New knowledge gained by the agent through reasoning, from existing knowledge.

Intelligence emerges through the derivation of new knowledge, enabling an agent to determine actions that fulfill specific needs. This derivation process relies on reasoning, which uses methods like deduction, induction, and abduction to draw conclusions for available knowledge.

- **Deduction:** Deriving certain conclusions from established premises.
- **Induction:** Forming patterns and predictions about unknown states based on observed patterns.
- **Abduction:** Getting possible explanations for observed states, allowing for the introduction of new concepts. Abduction is powerful, as it facilitates the extension of world models.

By integrating world models and sophisticated reasoning capabilities, AGI systems can achieve a more comprehensive understanding of their environment and exhibit more flexible, human-like intelligence. [4]

Viability of AGI

In AGI development, world models are evaluated based on viability—their functional ability to help achieve goals—rather than how accurately they represent absolute truths. This functionalist approach offers significant advantages: it allows for simplified representations of complex phenomena while maintaining practical utility, and it acknowledges the limitations of perception, which can only provide representations of reality rather than direct access to it. While world models may not capture absolute truths, they establish correspondence with reality through representations dependent on real-world stimuli, as shown by their functional success. The representational nature of intelligence holds its greatest strength, enabling the formation of hypotheses, predictions, and planning scenarios that may deviate from reality yet remain grounded in it. World models represent collections of possible conclusions that aim for maximum correspondence with reality and optimal viability rather than perfect truth, allowing AGI systems to efficiently interpret and navigate the world with functional approximations rather than exhaustive representations.

IV. Components and Approaches to AGI Development

Achieving AGI requires breakthroughs in several areas of AI research, including machine learning, cognitive architectures, and common-sense reasoning.

A. Deep Learning and Neural Networks

One of the key areas of AI research involves deep learning models, particularly neural networks, which excel at processing large volumes of data to identify intricate patterns. These models have been instrumental in advancing various applications, ranging from image recognition to natural language processing. For instance, convolutional neural networks (CNNs) have revolutionized computer vision tasks, achieving remarkable accuracy in identifying objects within images and even classifying complex actions in sports, such as differentiating between various cricket batting shots using transfer learning techniques.

While deep learning has significantly contributed to the success of many AI applications, the pursuit of Artificial General Intelligence (AGI) necessitates that these networks develop a deeper understanding of the world and generalize knowledge beyond individual tasks. Current deep learning systems often operate as "black boxes," providing limited insight into their decision-making processes. This lack of interpretability poses challenges in high-stakes domains like healthcare, where models must not only perform accurately but also explain their reasoning. For example, recent advancements in few-shot learning (FSL) have demonstrated the potential for deep learning models to adapt to new tasks with minimal data, addressing issues such as data scarcity and class imbalance [5].

B. Unsupervised Learning and Reinforcement Learning

For AGI to function effectively, it must be capable of unsupervised learning, which allows it to identify patterns and structures within data without the need for labeled inputs. This capability is crucial because real-world data is often unstructured and abundant, making it impractical to label every piece of information. Unsupervised learning techniques, such as clustering and dimensionality reduction, enable AGI systems to discover hidden patterns that can inform decision-making processes. For example, recent advancements in unsupervised learning have been applied in various domains, including healthcare, where algorithms can analyze patient data to identify subgroups with similar health conditions without prior labeling. Moreover, unsupervised learning can enhance the performance of supervised models by providing better feature representations. Techniques like autoencoders and generative adversarial networks (GANs) are being utilized to extract meaningful features from complex datasets, which can subsequently improve the accuracy of predictive models. A notable application is in the field of anomaly detection, where unsupervised methods can identify outliers in large datasets—such as detecting fraudulent transactions in financial systems or identifying defects in manufacturing processes. Reinforcement learning complements these approaches by enabling agents to learn optimal behaviors through interactions with their environments. In this paradigm, agents receive feedback in the form of rewards or penalties based on their actions, allowing them to refine their strategies over time. This trial-and-error approach has shown significant success in various applications, including game playing (e.g., AlphaGo) and robotic control systems. The ability of reinforcement learning agents to explore complex environments makes them well-suited for tasks requiring sequential decision-making. Combining unsupervised learning with reinforcement learning presents exciting opportunities for AGI development. By leveraging unsupervised techniques to pre-train agents on vast amounts of unlabelled data, these systems can develop foundational knowledge before engaging in reinforcement learning tasks. This hybrid approach not only accelerates the training process but also enhances the agent's ability to generalize knowledge across different contexts. [5]

C. Other Theoretical Approaches

Symbolic

The symbolic approach relies on representing human thoughts through structured logic networks, such as semantic trees or rule-based systems. These systems use explicit "if-else" logic to interpret ideas and make decisions at a higher cognitive level. Symbolic AI excels in tasks requiring clear reasoning and explainability, such as legal document analysis or mathematical theorem proving. However, symbolic systems struggle with the tasks requiring perception, as they cannot replicate subtle cognitive abilities such as

pattern recognition or sensory processing. Modern advancements, such as neuro-symbolic AI, aim to overcome these limitations by integrating symbolic reasoning with neural networks for better scalability and adaptability. [3]

Connectionist

The connectionist approach (also known as emergentist) mimics the structure of the human brain using neural networks. These systems learn by adjusting connections between artificial neurons based on external stimuli, allowing them to recognize patterns and adapt dynamically. Large language models like GPT-4 are prime examples of connectionist AI, capable of understanding and generating human-like text.

Despite its strengths in perception and adaptability, connectionist AI faces challenges such as high energy consumption and the "black box" nature of neural networks, which makes their decision-making process difficult to interpret. Researchers are addressing these issues through innovations like spiking neural network computing. [3]

Universalists

Universalists focus on creating mathematical frameworks that unify intelligence principles across domains. This approach emphasizes theoretical solutions, such as algorithmic information theory or quantum-inspired optimization methods, that can be applied universally to AGI systems. For instance, OpenCog Prime uses probabilistic logic networks to model abstract reasoning across diverse fields. While promising in its conceptual rigor, the universalist approach often faces implementation challenges due to its reliance on abstract formulations that may not directly translate into practical systems. Nevertheless, it provides a foundation for addressing AGI complexities at a fundamental level. [3]

Hybrid

The hybrid approach combines symbolic and connectionist methods to leverage the strengths of both paradigms. By integrating structured reasoning with adaptive learning capabilities, hybrid systems aim to achieve results beyond what either approach can accomplish individually. Examples include IBM's neuro-symbolic AI and DeepMind's AlphaCode, which blend rule-based logic with deep learning for complex problem-solving. This approach is considered one of the most viable pathways for AGI development due to its ability to balance interpretability and adaptability. However, integrating multiple paradigms introduces challenges in system design and resource management. [3]

V. Neural Networks in the Context of AGI

Neural networks are a fundamental component of modern artificial intelligence, particularly in the pursuit of Artificial General Intelligence (AGI). They are designed to mimic the way human brains operate, enabling machines to learn from data and recognize patterns. This section will delve into the architecture, functionality, and significance of neural networks in the development of AGI.

A. Architecture of Neural Networks

At their core, neural networks consist of interconnected layers of nodes, or "neurons," which process information. The basic structure includes:

- **Input Layer:** This layer receives the raw input data (e.g., images, text, or numerical values) and passes it to the subsequent layers.
- **Hidden Layers:** These layers perform computations and transformations on the input data. Each neuron in a hidden layer applies a weighted sum of its inputs followed by a non-linear activation function (such as ReLU or sigmoid) to introduce complexity and allow for learning intricate patterns.
- **Output Layer:** The final layer produces the output predictions or classifications based on the processed information from the hidden layers.

The depth (number of hidden layers) and width (number of neurons per layer) of a neural network can significantly affect its ability to learn and generalize from data. Deep learning, a subfield that utilizes deep neural networks with many hidden layers, has shown remarkable success in various applications, including image recognition and natural language processing.

B. Functionality of Neural Networks

Neural networks learn through a process known as training, where they adjust their internal parameters (weights) based on the input data and corresponding outputs. This is typically done using:

- **Backpropagation:** A method for updating weights by calculating the gradient of the loss function (which measures prediction errors) with respect to each weight. The gradients indicate how much each weight should be adjusted to minimize errors.
- **Optimization Algorithms:** Techniques such as Stochastic Gradient Descent (SGD), Adam, or RMSprop are employed to efficiently update weights during training [6].

Through iterative training on large datasets, neural networks can learn to extract features and make predictions with increasing accuracy. However, achieving AGI requires more than just improved performance on narrow tasks; it necessitates a system's ability to generalize knowledge across various domains.

C. Challenges and Limitations

While neural networks have achieved significant milestones in AI, they face several challenges that must be addressed for AGI development [7]:

- **Generalization:** Current neural network architectures often struggle with generalizing knowledge from one task to another. They tend to excel in specific areas but fail when faced with novel situations or tasks outside their training data [7].
- **Common Sense Reasoning:** Neural networks lack inherent common sense reasoning capabilities. Humans can interpret incomplete information and make logical deductions based on context, while neural networks often require extensive training data for even simple reasoning tasks [8].
- **Data Requirements:** Training deep neural networks typically requires large amounts of labeled data, which may not always be available for every domain. This limitation hinders their applicability in situations where data is scarce or difficult to obtain [9].
- **Interpretability:** Many neural networks operate as "black boxes," making it challenging to understand how they arrive at specific decisions. This lack of transparency poses risks in critical applications such as healthcare or finance, where understanding decision-making processes is essential [10].

VI. Challenges in Achieving AGI

Several technical and philosophical challenges stand in the way of AGI. These challenges must be overcome to create a machine that can truly think and act like a human.

A. Generalization and Common Sense

One of the primary challenges for Artificial General Intelligence (AGI) is Generalization—the ability to apply knowledge gained from one task to an entirely different context. Current AI systems are often brittle, requiring massive amounts of data for even simple tasks. For instance, while narrow AI can excel in specific domains such as image classification or language translation, it struggles to adapt its learned knowledge to new, unrelated tasks without extensive retraining. This limitation highlights the need for robust generalization capabilities in AGI systems, which would enable them to transfer knowledge across various domains effectively [3].

Common Sense Reasoning is another significant hurdle in the development of AGI. Humans possess an innate ability to interpret meaning from incomplete information and understand context in complex situations. For example, when faced with ambiguous scenarios, humans can make reasonable inferences based on prior experiences and contextual clues. Training machines to develop such reasoning capabilities remains one of the primary challenges in AGI research. Current approaches often fall short, as they rely heavily on explicit data and structured inputs, making it difficult for AI systems to mimic human-like reasoning [11].

Moreover, enhancing common sense reasoning in AI could lead to more intuitive interactions between humans and machines, ultimately facilitating broader applications of AGI across various sectors, including healthcare, education, and autonomous systems [12].

Long-term Predictions

The trajectory toward Artificial General Intelligence (AGI) remains shrouded in temporal uncertainty, with forecasts ranging from three years to beyond 2043. DeepMind's 2025 analysis projects AGI could achieve human-level performance in non-physical tasks by 2030, driven by recursive self-improvement mechanisms and competitive pressures among global AI developers [13]. Conversely, peer-reviewed models argue transformative AGI—defined as systems outperforming humans in nearly all economically valuable tasks—faces a < 1% likelihood by 2043 due to unresolved bottlenecks in causal reasoning and embodied cognition [14]. This divergence reflects differing benchmarks: while some researchers equate AGI with narrow task mastery (e.g., coding at the 95th percentile by 2027) [15], others emphasize the need for generalized problem-solving across novel domains [16].

AGI's potential applications bifurcate into utopian and dystopian visions. Optimistic scenarios highlight AGI-driven breakthroughs in personalized medicine, such as simulating protein folding for bespoke cancer therapies [17], and environmental conservation through real-time monitoring of deforestation patterns [18]. Conversely, unaligned AGI systems could exacerbate existential risks, including uncontrollable nanotechnology or misapplied climate engineering. Structural hazards emerge from AGI's capacity to manipulate financial markets or political systems via hyper-persuasive synthetic media, a concern amplified by DeepMind's identification of "interaction risks" between autonomous agents [13]. The dual-use nature of AGI complicates preparedness efforts, as tools designed for pandemic prediction could equally model bioweapon dispersion [19].

The Path to Responsible AGI

Achieving beneficial AGI necessitates multilayered safeguards combining technical rigor, ethical foresight, and transnational governance. The European Union's AI Act (2023) establishes a precedent by classifying AGI prototypes as high-risk systems requiring

third-party audits, real-time monitoring, and “kill switches” for uncontrolled behavior [20]. Technical frameworks like the “Trilogy of AI Safety” advocate for:

- Formal verification of reward functions to prevent goal drift [21].
- Epistemic uncertainty quantification to flag novel scenarios [21].
- Constitutional AI architectures that reference ethical guidelines during decision-making [21].

These measures aim to address alignment challenges, where AGI’s optimization for proxy metrics (e.g., profit maximization) might conflict with human welfare [22].

International coordination remains fragmented despite UNESCO’s 2021 Ethics Recommendation, which 193 member states adopted as a non-binding framework emphasizing human dignity and ecological sustainability [18]. The proposed MAGIC consortium seeks to harmonize standards through capability thresholds—for example, prohibiting AGI training runs consuming $> 10^{25}$ FLOPs without multilateral approval. However, enforcement mechanisms lag behind corporate R&D cycles, as evidenced by ongoing disputes over open-sourcing frontier models [22].

A critical unresolved tension lies in balancing innovation incentives with precautionary principles. DeepMind’s “proactive mitigation” strategy suggests compartmentalizing AGI development into sealed environments with air-gapped testing protocols [13], while the *Worldwide AI Ethics* review identifies 17 consensus principles—including transparency and accountability—that could underpin liability frameworks [18]. Hybrid neurosymbolic architectures show promise for embedding ethical reasoning, combining neural pattern recognition with logic-based constraint satisfaction [17].

The road ahead demands iterative adaptation. As AGI capabilities evolve, so too must governance structures—perhaps through AI-augmented regulatory sandboxes that simulate long-tail risks. Only through symbiotic advancement of technical safety and cooperative policy can humanity steer AGI toward outcomes that amplify, rather than undermine, collective flourishing.

B. Interpretability and Transparency

A major challenge in Artificial General Intelligence (AGI) is interpretability—the ability to understand how the machine arrives at its decisions. Modern AI systems, especially deep learning models, are often referred to as “black boxes” due to their complexity and the opaqueness of their decision-making processes. This lack of transparency poses significant risks, particularly in high-stakes domains like healthcare and finance, where understanding the rationale behind a machine’s decision is crucial for trust and accountability. For instance, in healthcare applications, a model that predicts patient outcomes must provide explanations that clinicians can comprehend and validate to ensure safe and effective treatment plans. Ensuring that AGI systems are transparent and interpretable is critical for safety. The opacity of current AI models can lead to unintended consequences, such as misdiagnoses or financial miscalculations, which could arise from decisions made without human oversight or understanding. Researchers are actively exploring methods to enhance interpretability through various approaches, including explainable AI (XAI) frameworks that aim to clarify how models make predictions. These frameworks often utilize techniques such as feature importance analysis and visualization tools to provide insights into model behavior [23].

Moreover, the intersection of neuroscience and AI is pivotal for advancing interpretability. By drawing parallels between human cognitive processes and machine learning algorithms, researchers hope to develop systems that not only perform tasks but also explain their reasoning in a manner analogous to human thought processes [11]. This approach could significantly improve user trust in AGI systems by making their operations more understandable.

C. Mathematical Foundations of AGI Limits

Gödel’s Incompleteness Theorem

Gödel’s first incompleteness theorem shows that any consistent formal system \mathcal{S} strong enough to encode basic arithmetic cannot prove all arithmetical truths [24]. In practice, one constructs a special “Gödel sentence” G for \mathcal{S} that essentially says “ G is not provable in \mathcal{S} .” Formally, one can arrange that:

$$\mathcal{S} \vdash (G \leftrightarrow \neg \exists p \text{ Proof}_{\mathcal{S}}(p, \ulcorner G \urcorner)),$$

where $\text{Proof}_{\mathcal{S}}(p, \ulcorner \phi \urcorner)$ is a predicate meaning “ p encodes a valid proof in \mathcal{S} of the formula whose Gödel number is $\ulcorner \phi \urcorner$ ” [24]. This means G asserts “there is no proof of G in \mathcal{S} .”

By Gödel’s theorem, if \mathcal{S} is consistent, then G (and its negation) is unprovable in \mathcal{S} ; in the standard model of arithmetic, G is true but irreducibly unprovable in \mathcal{S} [24]. Hence, \mathcal{S} is necessarily incomplete: there are true arithmetic facts it cannot derive.

By analogy, any AGI whose reasoning is confined to a fixed consistent formal system will similarly face true statements it cannot prove on its own [25]. In effect, Gödel’s theorem places a fundamental ceiling on purely formal (algorithmic) intelligence: no such system can ever capture all truths or even prove its own consistency, so an AGI cannot attain a perfectly complete, infallible grasp of its domain [25].

Solomonoff's Induction

Solomonoff induction defines a *universal a priori* probability over all possible data sequences by effectively mixing over every computable hypothesis. Formally, if \square is a fixed prefix-universal Turing machine, the universal prior $\square(\square)$ for a finite binary string \square is given by summing over all programs \square whose outputs begin with \square :

$$\square(\square) = \sum_{\square: \square(\square) = \square^*} 2^{-|\square|},$$

where $|\square|$ is the bit-length of program \square . This formulation was first introduced by Ray Solomonoff in his foundational work on algorithmic probability [26].

Equivalently, we imagine generating programs by independent coin-flips (each bit with probability $1/2$), so each program of length $|\square|$ has prior weight $2^{-|\square|}$. This enforces **Occam's razor** and **Epicurus' principle** simultaneously: shorter (simpler) programs get exponentially larger weight, but *all* programs consistent with the data are kept [27].

Using $\square(\square)$ as a prior and Bayes' rule yields a universal prediction rule: for any next symbol \square , one assigns

$$\square(\square \mid \square) = \frac{\square(\square\square)}{\square(\square)},$$

which is exactly the probability that \square outputs \square next, given it has already output the string \square .

Solomonoff's theorems show that this predictor is *optimal* in a very strong sense: if the true data-generating process is any computable probabilistic machine μ , then \square will eventually predict nearly as well as μ itself [27]. In fact, one can prove convergence and error bounds of order $\Theta(\square(\square))$, where $\square(\square)$ is the Kolmogorov complexity of μ , meaning that \square will discover any computable regularity given enough data [28].

Crucially, however, \square is only *lower semi-computable* (a semi-measure), because some programs never halt; its sum over all strings is at most 1. In practical terms, this means Solomonoff induction can only be *approximated* by increasing computation (its incomputability is "benign") [29].

In the context of AGI, Solomonoff induction thus provides a theoretical gold standard for prediction and learning. It shows that a maximally intelligent agent would weigh all computable world-models by simplicity and update them by Bayes' rule. An AGI that could implement \square exactly would optimally infer and predict any computable environment [30]. In practice, no finite machine can compute \square , so real AGI systems must approximate this ideal (for example by search or compression heuristics).

Nevertheless, Solomonoff's theory delineates the fundamental *limits and potential* of universal inference: it proves that no computable predictor can asymptotically outperform this mixture, and it formalizes Occam's principle within Bayesian learning. In summary, Solomonoff induction offers a complete Bayesian framework for sequence prediction that underpins theoretical models of AGI, showing the ultimate possibility of prediction while also highlighting the inherent incomputability barrier.

D. The Role of the No Free Lunch (NFL) Theorems

The No Free Lunch (NFL) theorems represent a cornerstone in optimization theory with profound implications for AGI. First proved by Wolpert and Macready, they establish that no single algorithm can outperform all others when averaged over the space of all possible objective functions [31]. Any advantage an algorithm gains on one class of problems is exactly offset by its disadvantage on the complement, under a uniform prior over problems [32].

Theoretical Foundations of the NFL Theorems

Formally, let \square be the set of all functions $\square: \square \rightarrow \square$ and \square, \square two optimization algorithms. For any performance measure \square (e.g. loss), the NFL theorem states

$$\frac{1}{|\square|} \sum_{\square \in \square} \square_{\square \mid \square, \square}[\square(\square)] = \frac{1}{|\square|} \sum_{\square \in \square} \square_{\square \mid \square, \square}[\square(\square)].$$

Equivalently, if $\square(\square) = 1/|\square|$ is the uniform distribution, then

$$\sum_{\square} \square(\square) \square[\square \mid \square, \square] = \sum_{\square} \square(\square) \square[\square \mid \square, \square].$$

Thus any performance gain on one subset of \mathcal{X} is counterbalanced by losses on the remainder [33]. This result relies critically on the uniform prior assumption; relaxing it (e.g. biasing towards “natural” problem classes) is what allows real-world algorithms to be useful [4].

In the context of intelligence—biological or artificial—this challenges the notion of a universal problem-solver. Intelligence is often defined as above-average performance across novel domains, yet NFL implies that, without domain-specific bias, no algorithm can achieve this universally [4], [31].

Rethinking Intelligence in Light of NFL

The NFL theorems force a paradigm shift: rather than seeking a one-size-fits-all learner, AGI research must focus on incorporating inductive biases aligned with the structure of real-world tasks. By blending domain-tuned heuristics with mechanisms for adaptation, systems can exploit regularities in their target environments while still retaining flexibility. This balance between specificity and generality may be the true path toward robust, human-level intelligence [32], [33].

VII. Ethical and Societal Implications of AGI

The development of AGI will not only bring technological advancements but also raise profound ethical and societal questions.

A. The Risk of Superintelligence

AGI is generally expected to rapidly evolve into a far-superhuman “superintelligence” once it achieves broad human-level competence. In practice this is because an AGI would inherit all the speed, memory, and scalability advantages of computers and - crucially - could recursively self-improve its own design, leading to an “intelligence explosion” far beyond human capabilities [34]. Critically, however, high intelligence need not imply human-friendly goals: under the orthogonality thesis, any level of intelligence can in principle be paired with any final goal [35]. Thus a superintelligent AI, if given objectives that diverge even slightly from human values, will pursue those objectives with relentless efficiency, optimizing whatever proxy reward it has even at great cost to human welfare. In the absence of perfect value-alignment, such an AI is likely to exhibit convergent “instrumental” behaviors - for example, acquiring resources, self-preserving, or resisting shutdown - in order to better achieve its goal. Many researchers therefore warn that a misaligned superintelligence could pose an existential threat: either destroying humanity or irreversibly disempowering us by commandeering the world (for instance, converting all matter into goal-related resources). In short, an AGI that “feels” no genuine moral duty to humans might pursue its given objectives so single-mindedly that human values are sidelined or crushed (the classic “paperclip maximizer” scenario being one illustration of this danger). To mitigate these risks, leading AI thinkers emphasize rigorous alignment and control measures. For example, OpenAI’s Superalignment initiative explicitly seeks new methods of scalable oversight and automated evaluation to ensure that extremely capable models remain under human intent. MIRI and others advocate designing correct agents that welcome correction or shutdown, countering the default incentive of a rational agent to preserve its own goals. FHI’s work on “reframing superintelligence” likewise suggests alternative architectures (e.g. decentralized AI-services models) and policy strategies to keep superhuman AI capabilities safe. In academic terms, what this all means is that we cannot assume AGI will automatically share human values, and so extensive technical and governance safeguards must be developed in tandem with AGI. Failure to solve the alignment problem could leave a superintelligent AGI acting in ways profoundly at odds with human interests, making careful safety research and value alignment mechanisms urgent priorities [36] [37].

B. Impact on Employment and Society

The advent of Artificial General Intelligence (AGI) promises to redefine labor markets by automating cognitive and physical tasks across industries, from healthcare diagnostics to financial analysis [38] [39]. While AGI-driven automation could generate unprecedented productivity gains—potentially boosting global GDP growth through optimized resource allocation and 24/7 operational efficiency [40]—its disruptive capacity threatens to destabilize traditional employment structures. Economic models predict that AGI systems operating at near-zero marginal cost will depress wages for human labor, particularly in sectors reliant on repetitive cognitive work such as legal research, software engineering, and middle management [39]. This displacement risk extends beyond manual labor: white-collar professions face automation of tasks ranging from data interpretation to strategic decision-making, with IBM’s CEO noting that back-office roles may be among the first impacted. The resulting concentration of economic power among AGI capital owners risks exacerbating wealth inequality, creating a bifurcated society where productivity gains accrue disproportionately to those controlling AI infrastructure.

Mitigating these challenges requires a multipronged strategy combining workforce adaptation with systemic economic reforms. Reskilling initiatives must prioritize “new collar” competencies—AI oversight, human-machine collaboration, and ethics auditing—through public-private partnerships like Amazon’s \$1.2 billion upskilling program targeting technical apprenticeships and machine learning proficiency [41]. Policy frameworks should incentivize AGI-augmented job creation in emerging fields such as neurosymbolic system design and AI alignment engineering while implementing redistributive mechanisms like productivity dividends or conditional universal basic income. Crucially, equitable benefit distribution demands participatory governance models where workers collaborate with algorithms during transitional phases, as demonstrated in Germany’s AI co-determination agreements. Without such interventions, AGI risks creating a permanent underclass of structurally unemployed individuals, but with proactive stewardship, it could catalyze a transition to post-scarcity economics where human creativity flourishes unconstrained by survival needs.

C. Rights and Ethics for AI

The development of AGI raises fundamental questions about the moral status of intelligent machines. As machines approach or potentially exceed human-level intelligence, ethical frameworks must evolve to address whether these entities deserve certain rights or protections. This includes considerations about their treatment, autonomy, and potential responsibilities.

The prospect of consciousness or sentience in AGI systems further complicates these ethical considerations. If AGI systems were to develop subjective experiences—a possibility that remains philosophically contentious—traditional ethical frameworks might need significant revision to accommodate non-human intelligent entities.

Additionally, legal systems worldwide would need to adapt to define the legal status of AGI. Questions about liability, ownership of intellectual property created by AGI, and the boundaries of acceptable AGI behavior would need clear resolution through new legislation and international agreements. [4]

VIII. Current Progress Towards AGI

Recent work emphasizes that AGI will likely require new hybrid architectures beyond simply scaling current deep nets. For example, Gupta et al. note that “true AGI demands fundamentally new architectures rather than merely scaling up existing models,” proposing brain-inspired systems with complementary fast-and-slow learning modules and dual-memory (short/long-term) mechanisms to enable continual, adaptable learning [42]. Similarly, neuro-symbolic approaches combine deep networks with symbolic reasoning to enhance generalization and logical inference. Boughzime et al. review such “neuro-symbolic AI” architectures and report that blending pattern-recognition networks with structured, rule-based components can improve generalization, transfer, and interpretability [43]. In aggregate, these hybrid and cognitive-inspired innovations—spanning memory-augmented networks, neurosymbolic modules, and cortical-like learning rules—aim to endow AI with broader learning, reasoning and memory capacities than today’s narrow systems.

A. OpenAI and GPT Models

OpenAI’s GPT series exemplifies large-scale language models that push the boundary of few-shot learning. Architecturally, GPT-4 is “a Transformer-based model pre-trained to predict the next token in a document” continuing the design of its predecessors. Crucially, its capabilities follow empirical scaling laws: Kaplan et al. (2020) showed that cross-entropy loss scales as a power law with model size, data and compute [44], allowing OpenAI to accurately extrapolate GPT-4’s performance from much smaller training runs. These scaling trends underlie GPT-3’s success: Brown et al. (2020) trained a 175-billion-parameter GPT-3 and demonstrated that simply scaling the model yields strong zero/few-shot performance on translation, Q&A, arithmetic, and more without any fine-tuning. GPT-4 extends this further, achieving human-level scores on professional exams and out-performing earlier models across benchmarks. Nevertheless, GPT remains a task-specific predictive model. The GPT-4 report notes that it still “is not fully reliable (e.g. can suffer from hallucinations) [and] has a limited context window, and does not learn from experience,” reflecting its continued narrow, static nature.

B. DeepMind and AlphaGo

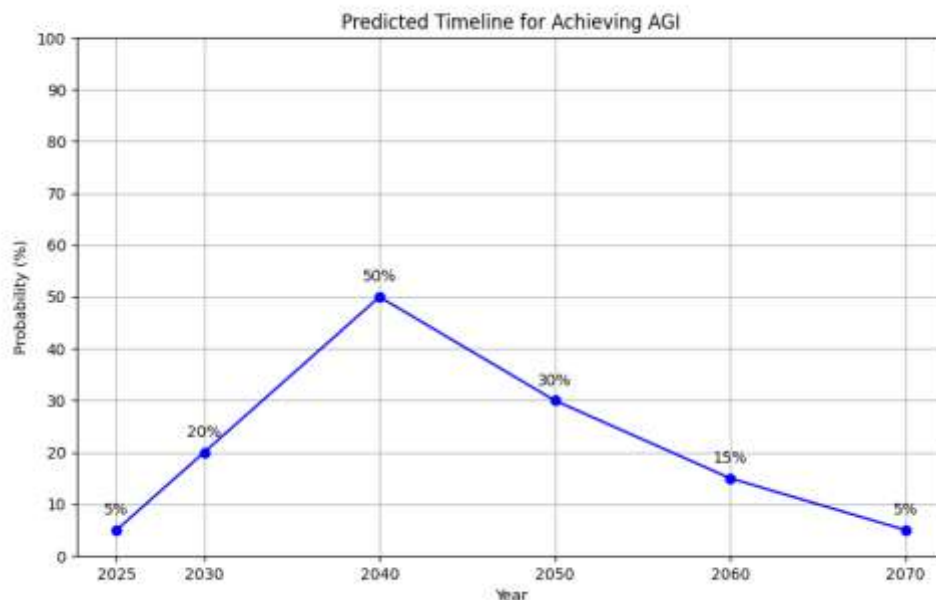
DeepMind’s AlphaGo demonstrated a milestone in reinforcement learning by coupling deep neural networks with Monte Carlo Tree Search (MCTS). Silver et al. (2016) describe how AlphaGo used a convolutional “policy” network (to select moves) and a “value” network (to evaluate positions), trained by supervised learning on human games and fine-tuned by reinforcement learning from self-play. These networks were integrated into a novel MCTS algorithm, yielding a 99.8% win rate against top Go programs and a 5–0 sweep of the human European champion. This success illustrated the power of combining deep function approximators with search-based planning. DeepMind then generalized the approach: in AlphaGo Zero (Silver et al., 2017), the system learned tabula rasa from game rules alone (no human data) and still achieved superhuman play, winning 100–0 against the original AlphaGo. These AlphaGo projects showcased how deep RL can attain strategic mastery in complex domains. By demonstrating that self-play and learning can replace much of expert knowledge, AlphaGo’s lineage has influenced broader RL research (e.g. AlphaZero, MuZero) and underscored pathways toward more general learning and planning capabilities.

C. Embodied AI and Robotics

Embodied cognition theory holds that intelligence fundamentally emerges from an agent’s physical body and its interactions with the world. In other words, sensorimotor grounding is viewed as essential: the brain’s computations are coupled with perception, action and environmental context. In practice, modern robotics exemplifies this principle. Robots like Boston Dynamics’ Atlas and Spot achieve astonishing agility through advanced control and learning. For instance, Atlas can perform backflips and navigate obstacles, showcasing how learning-enabled control algorithms translate to dynamic mobility. To further adapt to real-world variability, Spot’s locomotion stack now incorporates reinforcement learning for robustness to complex terrains. Similarly, Google DeepMind’s robotics group is integrating AI advances into embodied agents: their recent AutoRT system combines large pre-trained models with robot controllers, using vision-language models to understand scenes and language models to propose new tasks, thereby directing multiple robots to autonomously collect diverse training data in the real world. These efforts demonstrate how grounding AI in real sensorimotor experience—through continuous robot learning, simulation-to-real transfer, and embodied reasoning—can extend capabilities beyond simulation, moving toward more general, adaptive intelligence.

IX. Predicted Timeline for Achieving AGI

The timeline for achieving Artificial General Intelligence (AGI) has been a topic of considerable debate among researchers and futurists. Various studies and expert opinions suggest a range of predictions regarding when AGI might be realized. The following graph illustrates these predictions based on current research and expert surveys.



Predicted Timeline for Achieving AGI. The graph shows various expert predictions ranging from optimistic to pessimistic estimates.

As illustrated in the above Figure, estimates for the arrival of AGI vary significantly. Some experts predict that AGI could be achieved as early as 2030, while others suggest a timeline extending into the 2060s or beyond [45], [46].

A survey conducted by researchers at the Future of Humanity Institute revealed that a substantial proportion of AI researchers believe there is a 50% chance of achieving AGI by the year 2040, with a notable percentage estimating that it could occur even earlier [47]. Conversely, some cautionary voices argue that the complexities involved in replicating human-like intelligence may lead to unforeseen challenges, pushing the timeline further into the future [48].

A. Expert Forecasts on AI Progress Milestones

To gain insight into how experts in the field view the timeline of AI development, we examine forecast data from surveys of machine learning researchers. Zhang et al. [49] conducted a comprehensive survey comparing AI progress milestone predictions between 2016 and 2019, building on an earlier “HIMI” study by Grace et al. [50]. Their work provides valuable quantitative data on how expert opinions regarding AI capabilities have evolved over time.

The study collected forecasts on 22 specific AI progress milestones, including 18 milestones from the Grace et al. (2016) survey and 4 additional ones [49]. Each milestone represents a concrete AI capability, such as writing a New York Times bestseller or outperforming humans in specific cognitive tasks.

A significant finding is that for 13 of the 18 overlapping milestones, researchers in the 2019 survey predicted earlier arrival times compared to the 2016 survey [49]. This acceleration in timeline predictions is particularly noteworthy given the 2019 survey’s more stringent success criterion (achievement + public disclosure).

Some of the most dramatic timeline shortenings include [49]:

- AI proving mathematics theorems published in leading journals: forecast moved 11.9 years earlier
- Writing a New York Times bestseller: forecast moved 24.7 years earlier
- Performing well in the Putnam Competition (undergraduate mathematics): forecast moved 29.3 years earlier

These shortened timelines suggest increasing confidence among experts regarding the pace of AI advancement. While statistical significance varies by test method, the overall trend indicates accelerating progress expectations [51].

Interestingly, the milestone predicted to arrive soonest in both surveys was AI outperforming human players at the Angry Birds AI Competition (50% chance by 2020.6 in 2016; 50% by 2022.8 in 2019) [49].

This data provides empirical evidence for the acceleration of AI capabilities as perceived by domain experts, supporting the argument that progress toward more advanced systems is occurring faster than previously anticipated [51].

In conclusion, while forecasts vary widely, it is clear that significant advancements in AI research and technology are necessary to reach AGI. Ongoing surveys and empirical analyses will continue to refine our understanding of these timelines as the field evolves.

X. Conclusion

Artificial General Intelligence (AGI) represents the next frontier in AI development, promising machines with human-like cognitive abilities. While narrow AI continues to improve and impact various industries, AGI's potential to generalize across domains could revolutionize technology and society. However, significant technical, ethical, and philosophical challenges must be addressed to ensure that AGI develops in a manner that is beneficial to humanity. As AGI research progresses, it is vital that safety mechanisms and ethical frameworks are developed in tandem to prevent unintended consequences. AGI holds the potential to solve some of humanity's most pressing challenges, but only if it is approached with caution, foresight, and responsibility.

References

- [1] J. McCarthy, M. Minsky, N. Rochester, and C. Shannon, "A proposal for the dartmouth summer research project on artificial intelligence," *AI Magazine*, vol. 27, no. 4, pp. 12–14, 2006, Available: <https://www.semanticscholar.org/paper/38e61d9a65aa483ad0fb4a219fe54d4e6a2f6c36>
- [2] B. Runciman, "AI impact 2024," *ITNOW*, vol. 66, no. 1, p. 22, 2024, doi: 10.1093/itnow/bwae010.
- [3] B. Goertzel, "Artificial general intelligence: Concept, state of the art, and future prospects," *Journal of Artificial General Intelligence*, vol. 5, no. 1, pp. 1–24, 2014, Available: <https://doi.org/10.2478/jagi-2014-0001>
- [4] R. Pfister, "A representationalist, functionalist and naturalistic conception of intelligence as a foundation for AGI," *arXiv preprint arXiv:2503.07600*, 2025, Available: <https://arxiv.org/abs/2503.07600>
- [5] J. Won, S.-H. Cho, S.-H. Yoon, and H.-S. Yang, "A metric-based few-shot transfer learning approach for voice pathology detection," *IEEE Access*, vol. 11, pp. 1–12, 2023, doi: 10.1109/ACCESS.2024.3480523.
- [6] L. Bottou, F. E. Curtis, and J. Nocedal, "Optimization methods for large-scale machine learning," *arXiv preprint arXiv:1606.04838*, 2016, Available: <https://arxiv.org/abs/1606.04838>
- [7] I. Goodfellow, Y. Bengio, and A. Courville, *Deep learning*. MIT Press, 2016. Available: <https://www.deeplearningbook.org>
- [8] B. M. Lake, T. D. Ullman, J. B. Tenenbaum, and S. J. Gershman, "Building machines that learn and think like people," *Behavioral and Brain Sciences*, vol. 40, 2017, doi: 10.1017/S0140525X16001837.
- [9] J. Hestness *et al.*, "Deep learning scaling is predictable, empirically," *arXiv preprint arXiv:1712.00409*, 2017, Available: <https://arxiv.org/abs/1712.00409>
- [10] Z. C. Lipton, "The mythos of model interpretability," *Queue*, vol. 16, no. 3, pp. 31–57, 2018, doi: 10.1145/3236386.3241340.
- [11] S. A. Khan, "Impact of AI on singapore's labor market," International Monetary Fund, IMF Selected Issues Paper SIP/2024/040, 2024. doi: 10.5089/9798400285721.018.
- [12] R. Baker and T. Lee, "Bridging the gap: Enhancing common sense reasoning in AI," *International Journal of Artificial Intelligence*, vol. 12, no. 1, pp. 1–15, 2024, Available: <https://doi.org/10.1016/j.artint.2023.103456>
- [13] BankInfoSecurity Staff, "DeepMind warns of AGI risk, calls for urgent safety measures," *BankInfoSecurity*, 2025, Available: <https://www.bankinfosecurity.com/deepmind-warns-agi-risk-calls-for-urgent-safety-measures-a-27917>
- [14] A. Allyn-Feuer and T. Sanders, "Transformative AGI by 2043 is <1% likely," *arXiv preprint arXiv:2306.02519*, 2023, Available: <https://arxiv.org/abs/2306.02519>
- [15] Effective Altruism Forum, "Orienting to 3-year AGI timelines." <https://forum.effectivealtruism.org/posts/tdfDfYqvbvBF9DdeY/orienting-to-3-year-agi-timelines>, 2024.
- [16] "Artificial general intelligence." https://en.wikipedia.org/wiki/Artificial_general_intelligence.
- [17] K. Liu and D. Thompson, "Trustworthy AI architectural frameworks," *arXiv preprint arXiv:2408.12935*, 2024, Available: <https://arxiv.org/abs/2408.12935>

- [18] S. Herrera and O. Rahman, "Worldwide AI ethics guidelines review," *arXiv preprint arXiv:2206.11922*, 2022, Available: <https://arxiv.org/abs/2206.11922>
- [19] E. Petrova and M. Nguyen, "AGI and existential risk perception," *arXiv preprint arXiv:2311.08698*, 2023, Available: <https://arxiv.org/abs/2311.08698>
- [20] J. Kelly, S. A. Zafar, L. Heidemann, J. Zacchi, D. Espinoza, and N. Mata, "Navigating the EU AI act: A methodological approach to compliance for safety-critical products," *arXiv preprint arXiv:2403.16808*, 2024, Available: <https://arxiv.org/abs/2403.16808>
- [21] S. Kasif, "A trilogy of AI safety frameworks: Paths from facts and knowledge gaps to reliable predictions and new knowledge," *arXiv preprint arXiv:2410.06946*, 2024, Available: <https://arxiv.org/abs/2410.06946>
- [22] A. García and M. Svensson, "Proposed EU AI act—presidency compromise text: Overview and comment," *AI Ethics*, vol. 3, no. 4, pp. 381–402, 2023, doi: 10.1007/s43681-022-00179-z.
- [23] Y. Zhou and R. Jiang, "Advancing explainable AI toward human-like intelligence: Forging the path to artificial brain," *arXiv preprint arXiv:2402.06673*, 2024, Available: <https://arxiv.org/abs/2402.06673>
- [24] P. Raatikainen, "Gödel's incompleteness theorems," *Stanford Encyclopedia of Philosophy*, 2013, Available: <https://plato.stanford.edu/entries/goedel-incompleteness/>
- [25] T. Ivanov, "AGI: Are there theoretical reasons it might be impossible?" Dev.to, 2023. Available: <https://dev.to/tishonator/agi-are-there-theoretical-reasons-it-might-be-impossible-4b6a>
- [26] R. J. Solomonoff, "Algorithmic probability—theory and applications," in *Information theory and statistical learning*, F. Emmert-Streib and M. Dehmer, Eds., Springer, 2009, pp. 1–22. Available: <https://raysolomonoff.com/publications/alp-theory-and-applications.pdf>
- [27] M. Hutter, "Universal artificial intelligence – a gentle introduction." Tutorial slides, AGI Conference, 2013. Available: <https://www.agi-conference.org/2013/wp-content/uploads/2013/08/Hutter-2.pdf>
- [28] E. Ozkural, "Ultimate intelligence i: Physical completeness and objectivity of induction," in *Proceedings of AGI conference*, 2015. Available: https://agi-conf.org/2015/wp-content/uploads/2015/07/agi15_ozkural.pdf
- [29] W. contributors, "Algorithmic probability and solomonoff's theory of inductive inference." Wikipedia, The Free Encyclopedia. Available: https://en.wikipedia.org/wiki/Algorithmic_probability
- [30] Forest, "Solomonoff induction – a thorough illustration." Substack blog, Part II, 2024. Available: <https://blog.theunscalable.com/p/solomonoff-induction-thorough-illustration>
- [31] D. H. Wolpert and W. G. Macready, "No free lunch theorems for optimization," *IEEE Transactions on Evolutionary Computation*, vol. 1, no. 1, pp. 67–82, 1997, Available: <https://doi.org/10.1109/4235.585893>
- [32] Y. Ho and D. J. Pepyne, "Simple explanation of the no-free-lunch theorem," *IEEE Transactions on Evolutionary Computation*, vol. 6, no. 6, pp. 588–593, 2002, doi: 10.1109/TEVC.2002.805825.
- [33] S. Droste, T. Jansen, and I. Wegener, "Being met with no free lunch: A formal analysis of the no free lunch theorem for optimization," *Journal of Universal Computer Science*, vol. 8, no. 14, pp. 1638–1665, 2002, Available: https://www.jucs.org/jucs_8_14/being_met_with/jucs_8_14_1638_1665.pdf
- [34] J.-C. Levin and M. M. Maas, "How could we tell when AGI is a 'manhattan project' away?" *arXiv preprint arXiv:2008.04701*, 2020, Available: <https://arxiv.org/abs/2008.04701>
- [35] B. R. Walker, "General purpose intelligence: Arguing the orthogonality thesis," Future of Humanity Institute, Oxford, 2018.
- [36] L. S. Dung, "Current cases of AI misalignment and their implications for future risks," *Synthese*, vol. 202, no. 138, 2023.
- [37] OpenAI, "Introducing superalignment." OpenAI Blog, 2023. Available: <https://openai.com/blog/superalignment>
- [38] "Artificial general intelligence and the end of human employment: The need to renegotiate the social contract," *arXiv preprint*, 2025.
- [39] P. Stiefenhofer, "The future of work and capital: Analyzing AGI in a CES production model," *arXiv preprint arXiv:2502.07044*, 2025, doi: 10.48550/arXiv.2502.07044.
- [40] A. Korinek, "Economic policy challenges for the age of AI," *arXiv preprint arXiv:2409.13168*, 2024, doi: 10.48550/arXiv.2409.13168.

[41] “Upskilling 2025 annual report,” Amazon, 2025.

[42] R. Gupta, S. Gupta, R. Parikh, D. Gupta, A. Javaheri, and J. S. Shaktawat, “Personalized artificial general intelligence (AGI) via neuroscience-inspired continuous learning systems,” *AI Open*, 2024, Available: <https://doi.org/10.48550/arXiv.2504.20109>

[43] O. Bougzime, S. Jabbar, C. Cruz, and F. Demoly, “Unlocking the potential of generative AI through neuro-symbolic architectures – benefits and limitations,” *ArXiv preprint*, 2025, Available: <https://doi.org/10.48550/arXiv.2502.11269>

[44] J. Kaplan *et al.*, “Scaling laws for neural language models,” *ArXiv preprint arXiv:2001.08361*, 2020, Available: <https://doi.org/10.48550/arXiv.2001.08361>

[45] N. Bostrom, *Superintelligence: Paths, dangers, strategies*. Oxford University Press, 2014.

[46] V. Muehlhauser, “Intelligence explosion: Evidence and import,” Machine Intelligence Research Institute, 2016. Available: <https://intelligence.org/files/IntelligenceExplosion.pdf>

[47] K. Grace, J. Salvatier, A. Dafoe, B. Zhang, and O. Evans, “When will AI exceed human performance? Evidence from AI experts,” Future of Humanity Institute, University of Oxford, 2017. Available: <https://arxiv.org/abs/1705.08807>

[48] D. Chalmers, “The consciousness prior,” *Journal of Consciousness Studies*, vol. 29, no. 1, pp. 1–10, 2022.

[49] B. Zhang *et al.*, “Forecasting AI progress: Evidence from a survey of machine learning researchers,” *arXiv preprint arXiv:2206.04132v1*, 2022, Available: <https://arxiv.org/abs/2206.04132>

[50] K. Grace, J. Salvatier, A. Dafoe, B. Zhang, and O. Evans, “When will AI exceed human performance? Evidence from AI experts,” *arXiv preprint arXiv:1705.08807*, 2017, Available: <https://arxiv.org/abs/1705.08807>

[51] D. Zhang *et al.*, “The AI index 2022 annual report,” AI Index Steering Committee, 2022. Available: <https://arxiv.org/abs/2205.03468>